

Big data architectures for Machine Learning and Data Mining

Apache Spark

Sagar Nagaraj Simha

Matrikel Number - 120797, sagar.nagaraj.simha@uni-weimar.de

Master's Computer Science for Digital Media - SS2019

Bauhaus University Weimar

What is Apache Spark?

A unified computing engine and a set of libraries for parallel data processing on computer clusters

Origin



Original author(s)	Matei Zaharia
Developer(s)	Apache Software Foundation, UC Berkeley AMPLab, Databricks
Initial release	May 26, 2014; 5 years ago ↗
Stable release	v2.4.3 / May 8, 2019; 37 days ago
Repository	https://github.com/apache/spark ↗
Written in	Scala, Java, Python, R ^[1]
Operating system	Microsoft Windows, macOS, Linux
Available in	Scala, Java, SQL, Python, R
Type	Data analytics, machine learning algorithms
License	Apache License 2.0
Website	spark.apache.org ↗

Why develop Apache Spark?

The Overall MapReduce Word Count Process

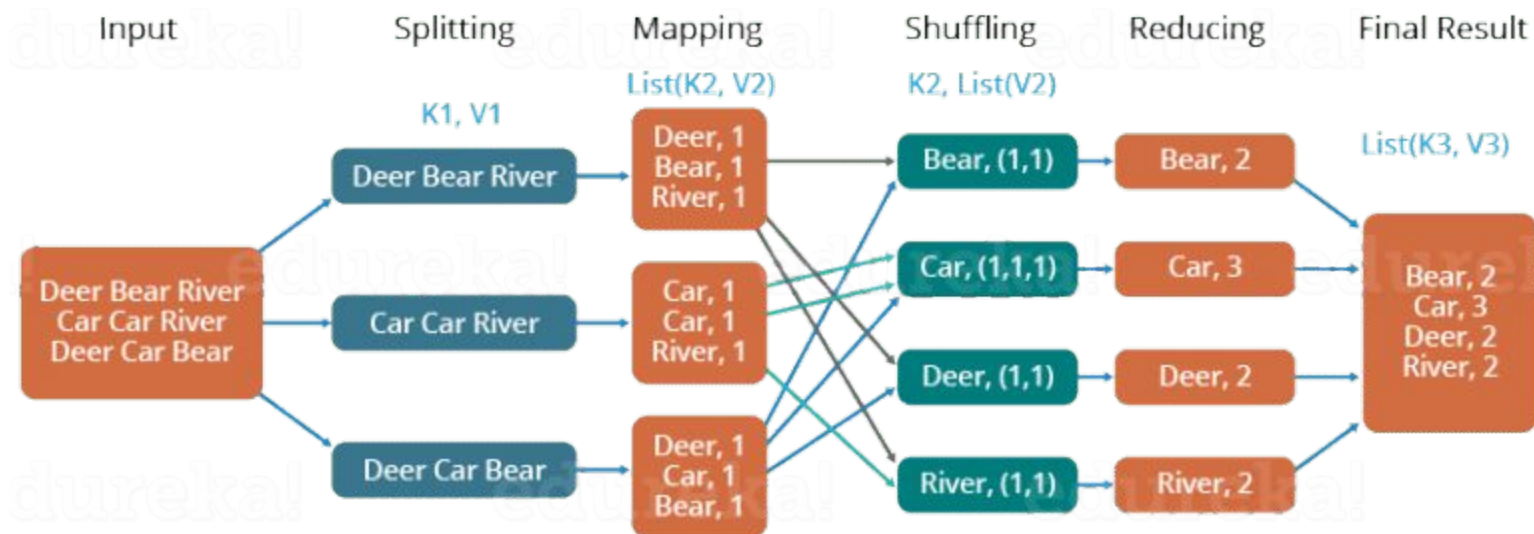




Image courtesy [7]

Apache Hadoop-MapReduce



**Data moving over
network in
Hadoop is slow**

Slow due to replication and disk I/O,
but necessary for fault tolerance

Apache Hadoop-MapReduce

MapReduce is inefficient for *multi-pass* and low latency requirement applications

- *Iterative algorithms*
- *Interactive data mining*
- *Streaming applications*

Hadoop included the Hadoop file system and MapReduce.

- Hard to run one of the systems without the other
- Incompatible with other cloud storage services

Challenges

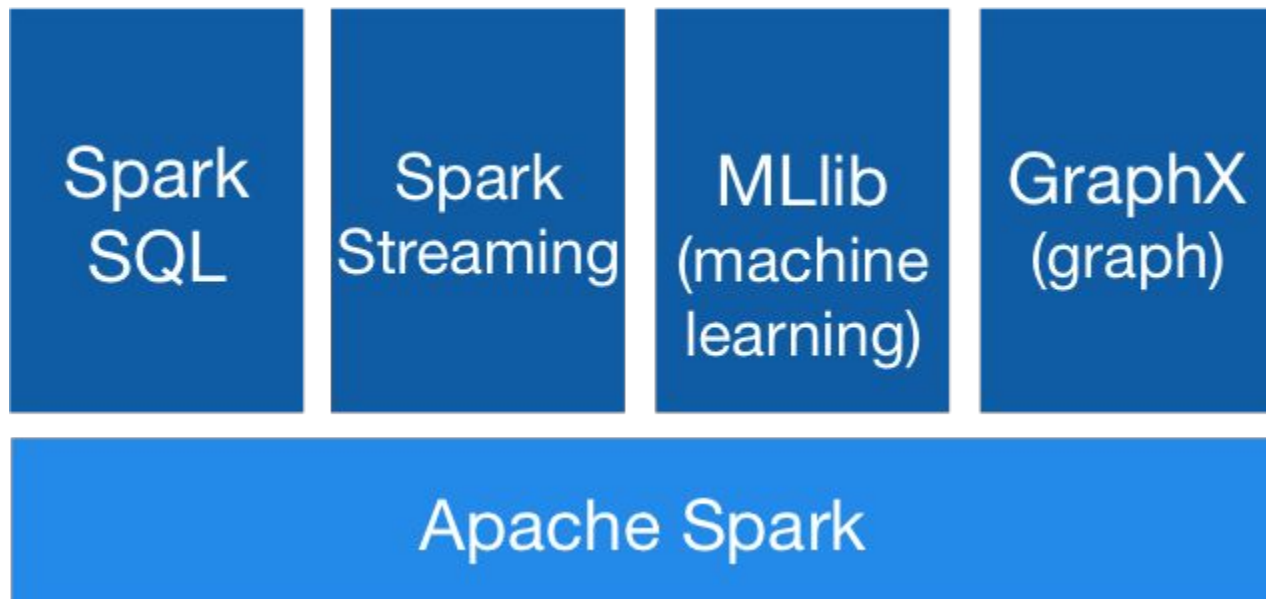
- Incorporate the previously mentioned features
- Have a distributed memory abstraction
- Fault tolerant and efficient

Apache Spark's Philosophy

“A new engine and programming model for data analytics”



1. Unified
 - a. Combines batch, interactive and streaming (incremental processing)
2. Computing Engine
 - a. Supports HDFS, Cassandra, Kafka etc and focuses on Computing - Moving data is expensive.
3. Libraries
 - a. Spark SQL - SQL and structured data
 - b. MLlib - Machine learning
 - c. GraphX - Graph analytics
 - d. Spark Streaming and the newer Structured Streaming - Stream processing
 - e. Many others on <https://spark-packages.org/> (Ex: Connectors to cloud storages, ML algorithms etc)
4. Resilient Distributed Datasets at the Lowest level for efficient fault tolerance

The Stack




Resilient Distributed Datasets (RDDs)

Apache Hadoop-MapReduce



Slow due to replication and disk I/O,
but necessary for fault tolerance

Spark - In Memory Data Sharing using RDDs



Fault Recovery

RDDs track the graph of transformations that built them (their lineage - map, filter, join etc...) to rebuild lost data

- Each row in the RDD input is a partition stored in a different machine in the cluster
- The transformations (in green box) are applied to each partition
- A batch data (blue box) may be recovered by running the corresponding transformation that produced it

RDD input

```
w1 w2 w3  
w1 w1 w3  
w1 w3 w3
```

RDD words

```
[w1, w2, w3, w1, w1, w3, w1, w3, w3]
```

RDD words with initial counts

```
[(w1,1), (w2,1), (w3,1), (w1,1), w1,1),(w1,1),  
(w3,1), (w1,1), (w3,1), (w3,1)]
```

RDD words with final counts



```
[(w1,4), (w2,1), (w3,4)]
```

flatMap(lambda x: x.split(' '))



map(lambda x: (x,1))

reduceByKey(lambda x,y:x+y)

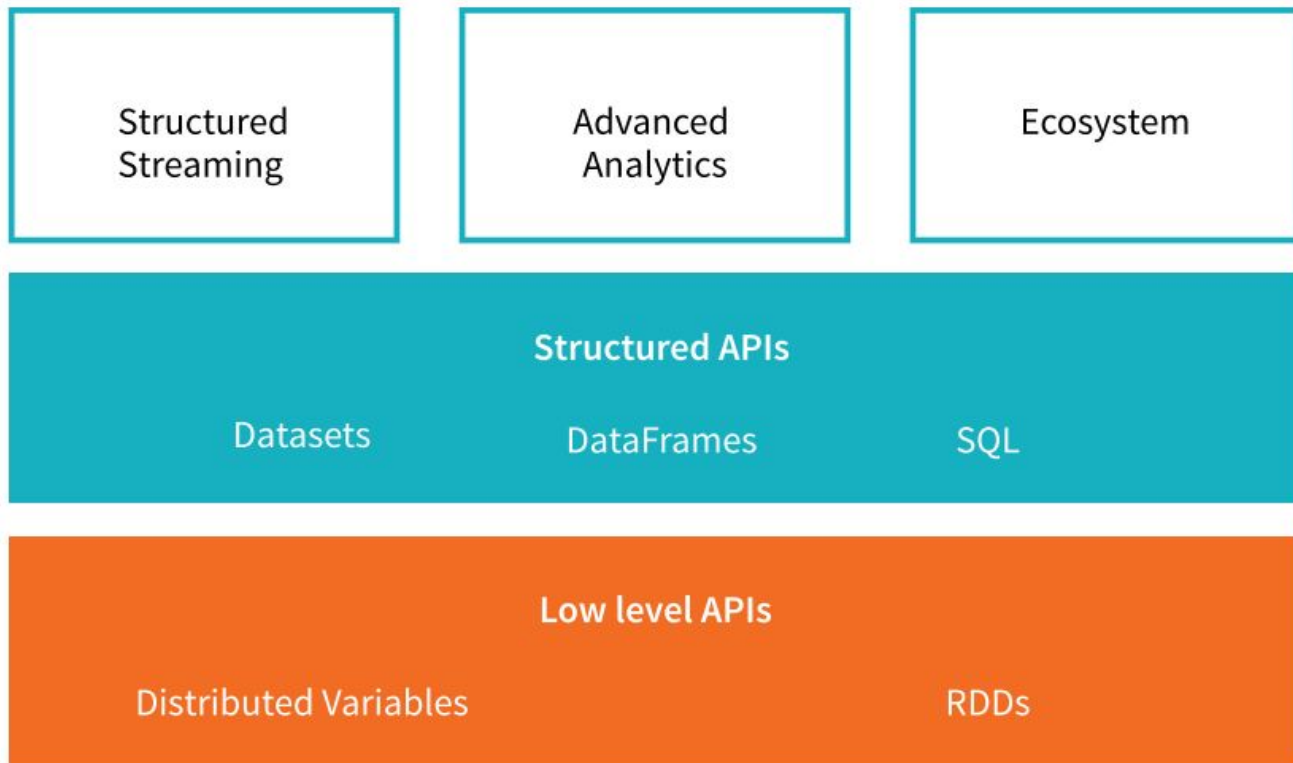
RDD Recovery - Fault Tolerance



RDD Recovery - Fault Tolerance

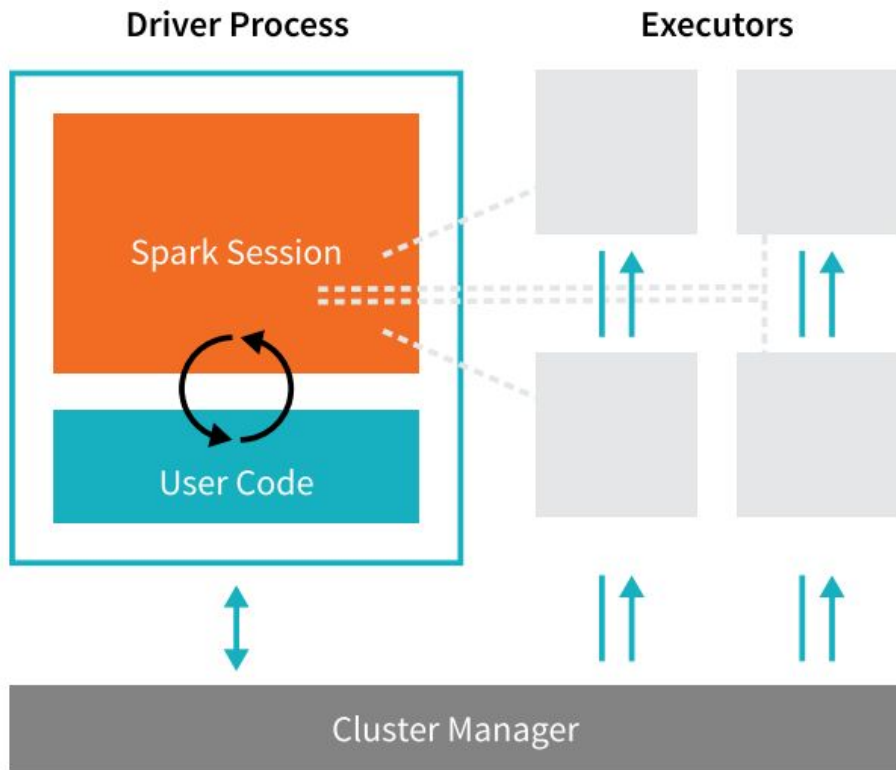


What Spark offers

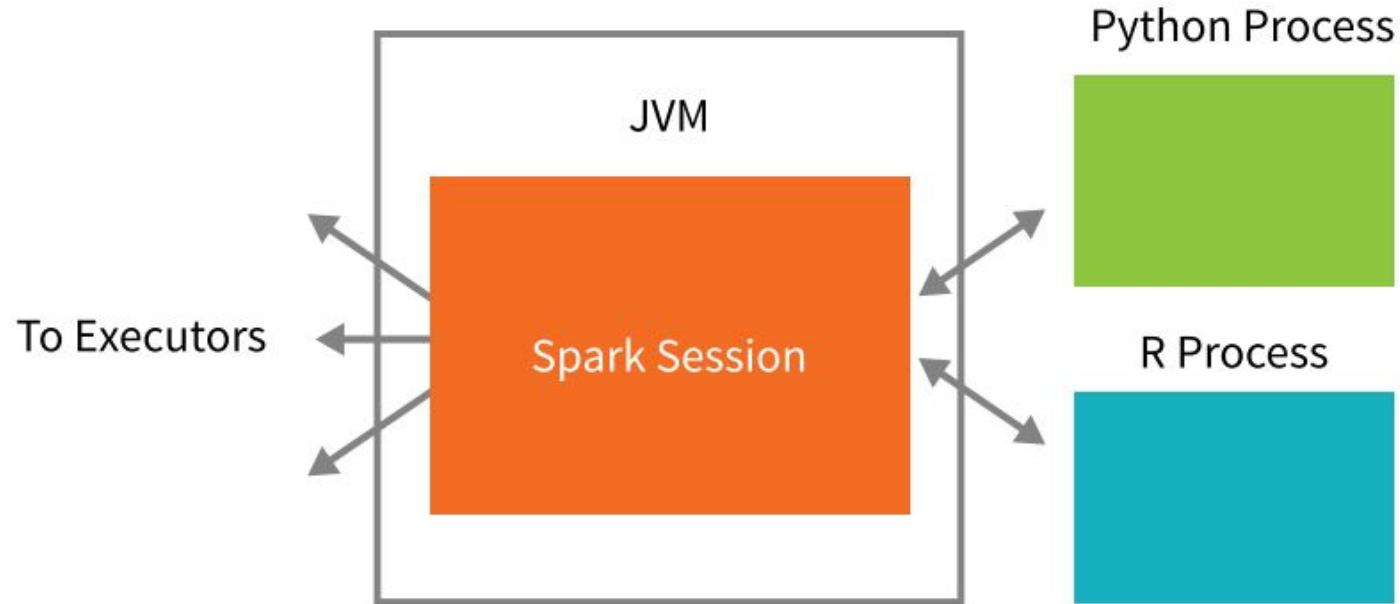


Architecture and workflow

Apache Spark architecture and workflow



Scala, Python & R hooks to launch a Spark Session



Spark Example: Log Mining

Load error messages from a log into memory and run interactive queries

```
lines = spark.textFile("hdfs://...")
```

base RDD

```
errors = lines.filter(startsWith("ERROR"))
```

transformation


```
errors.persist()
```

```
errors.filter("404" in errors).count()
```

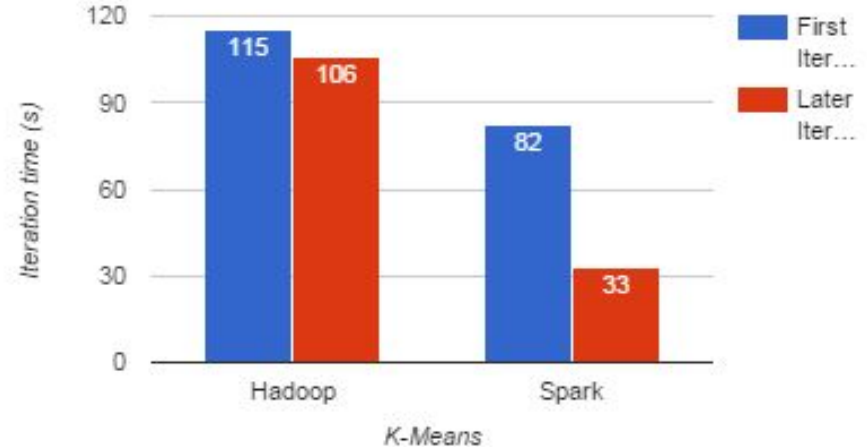

action!

```
errors.filter("405" in errors).count()
```

Result: full-text search on 1TB data on 100 machines
in 5-7sec vs. 170sec with on-disk data!



Performance - 10 iterations on 100GB data using 25-100 machines



Features/Advantages

1. Open-source
2. Speed - 100x faster than Hadoop for large scale data processing
3. Automatic fault tolerance
4. Unified Engine
5. Usability with multiple languages
6. Lazy Evaluation - “predicate pushdown”
7. Compatibility with other ecosystems
8. Supports interactive and production applications
9. Easy-to-use APIs for operating on large datasets

Companies/Products that use Spark

- Amazon
- Uber
- Baidu
- eBay Inc.
 - Using Spark core for log transaction aggregation and analytics
- Yandex
- PanTera
 - PanTera is a tool for exploring large datasets. It uses Spark to create XY and geographic scatterplots from millions to billions of datapoints.
- Vistar Media
 - Location technology company enabling brands to reach on-the-go consumers

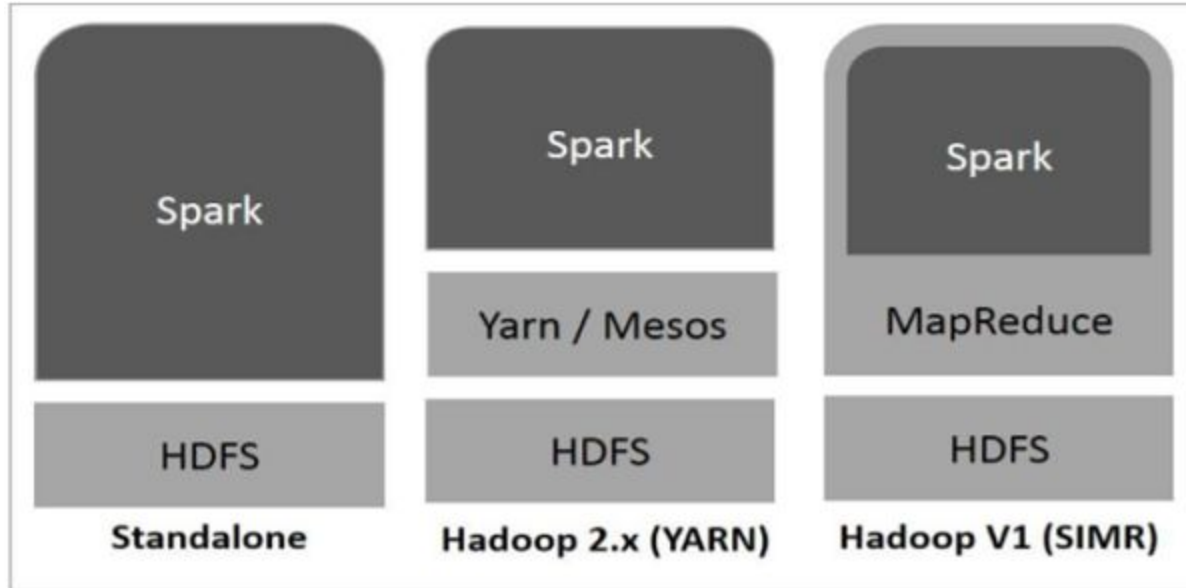
Research institutions that use Spark

- CERN Open Lab
 - From Collision to Discovery: Physics Analysis with Apache Spark
- European Gravitational Observatory
 - Machine Learning for Gravitational Wave signals classification in LIGO and Virgo
- Freeman Lab at HHMI
 - We are using Spark for analyzing and visualizing patterns in large-scale recordings of brain activity in real time
- NASA JPL - Deep Space Network
- Stanford DAWN
 - Research lab on infrastructure for usable machine learning, with multiple research projects that run over or accelerate Apache Spark.

Installation

For the Demo, we will run a Hadoop Compatible Apache Spark on a Vagrant VM machine with a driver(master) node and N(2) worker nodes

Apache Spark with other systems



We will use Apache Spark Standalone Cluster manager for demo

Apache Spark with other systems

Over Hadoop YARN Cluster,

<https://www.linode.com/docs/databases/hadoop/install-configure-run-spark-on-top-of-hadoop-yarn-cluster/>

<https://databricks.com/blog/2014/01/21/spark-and-hadoop.html>

<http://www.datumly.com/2017/08/apache-spark-2-2-in-a-virtual-machine-simple-getting-started-guide-to-run-spark-on-your-laptop/>

<https://medium.com/explore-artificial-intelligence/downloading-spark-and-getting-started-with-python-notebooks-jupyter-locally-on-a-single-computer-98a76236f8c1>

With Hadoop:

<https://www.davidadrian.cc/posts/2017/08/how-to-spark-cluster/>

With Jupyter Notebook:

<https://opensource.com/article/18/11/pyspark-jupyter-notebook>

DEMOS

Demo plan

1. Launching the VM cluster with Spark
2. Spark UI
3. Spark Shell (PySpark)
4. Features of Spark - RDDs, DataFrames, Lazy Evaluation
5. Flight Data Analysis
6. Word Count
7. Production Application - Calculation of Pi

Disadvantages

- No File Management System (Reliant on other storage)
- In-memory computations are expensive
- Spark MLlib has limited algorithms
- Manual Optimization
- Lower latency compared to Apache Flink

Research in Spark

<https://spark.apache.org/research.html>

Spark Summit - <https://databricks.com/sparkaisummit>

To deep dive further


- A-Gentle-Introduction-to-Apache-Spark (Databricks)
- Spark - The Definitive Guide - Big data processing made simple (O'Reilly)
- Learning Spark, Lightning-Fast Big Data Analysis (O'Reilly)

Dataset Used for Demo

<https://github.com/databricks/Spark-The-Definitive-Guide/tree/master/data>

References

- [1] <https://www.youtube.com/watch?v=dXG4yC8ICEI>
- [2] <https://spark.apache.org/research.html>
- [3] <https://www.whizlabs.com/blog/apache-spark-limitations/>
- [4] <https://spark.apache.org/powered-by.html>
- [5] <https://www.youtube.com/watch?v=L029ZNBG7bk&feature=youtu.be>
- [6] A Gentle Introduction to Apache Spark - By Databricks
- [7] <https://www.edureka.co/blog/mapreduce-tutorial/>



"After careful consideration of all 437 charts, graphs, and metrics,
I've decided to throw up my hands, hit the liquor store,
and get snockered. Who's with me?!"

Vielen Dank!